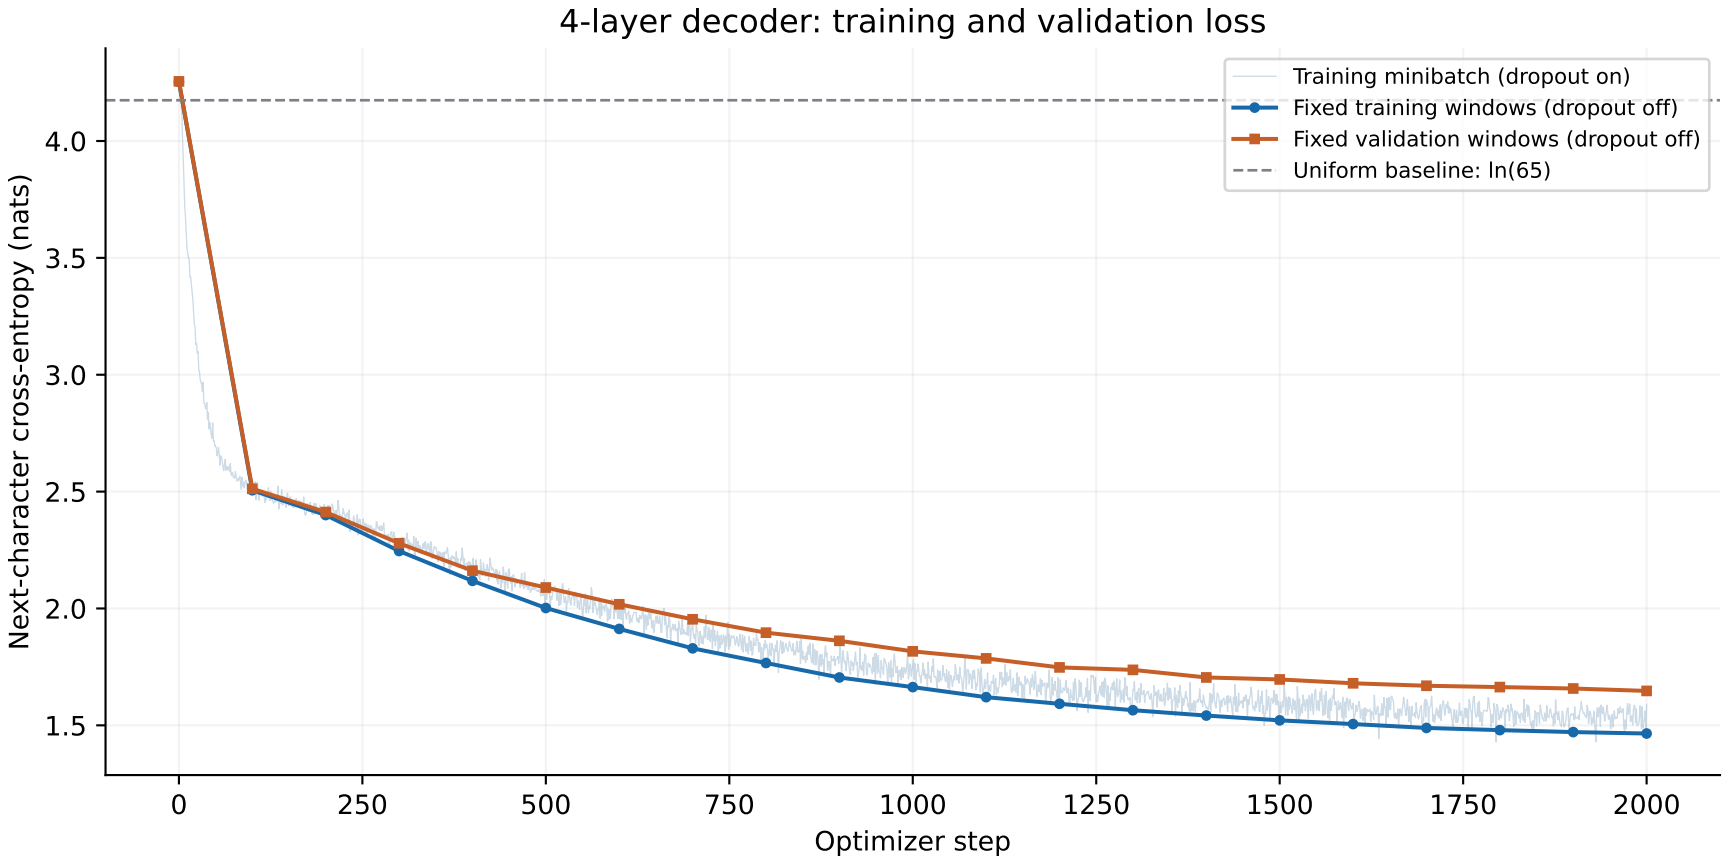
4-layer decoder: training and validation loss
Training minibatch (dropout on)
Fixed training windows (dropout off)
Fixed validation windows (dropout off)
Uniform baseline: ln(65)
Next-character cross-entropy (nats)
4.0
3.5
3.0
2.5
2.0
1.5
0
250
500
750
1000
1250
1500
1750
2000
Optimizer step